

Exercise Book

Ugo Dal Lago

Francesco Gavazzo

1 Mathematical Preliminaries

Exercise 1.1. What is the cardinality of the set S^n , where S is any finite alphabet? Prove your claim.

Exercise 1.2. What is the cardinality of the subset of $\{0, 1\}^{2n}$ consisting of all and only the *palindrome* words? Prove your claim.

Exercise 1.3. What is the cardinality of the subset of $\{0, 1\}^n$ consisting of all and only the words which have even *parity* (and the parity of a binary string is the number of occurrences of the symbol 1 inside it)? Prove your claim.

Exercise 1.4. Relate the following pair of functions (f_i, g_i) by way of $O(\cdot)$, $\Omega(\cdot)$ or $\Theta(\cdot)$ notation:

$$f_1(n) = n^2$$

$$g_1(n) = 4n^1 + 100 \log(n)$$

$$f_2(n) = n \log(n)$$

$$g_2(n) = 10n \log(\log(n))$$

$$f_3(n) = 2^n n^2$$

$$g_3(n) = 3^n$$

$$f_4(n) = 100n$$

$$g_4(n) = \frac{1}{100} n \log(n)$$

$$f_5(n) = \log^3(n)$$

$$g_5(n) = n^{\frac{2}{3}}$$

$$f_6(n) = 10^{-3} n^3$$

$$g_6(n) = 10^4 n^3 + 10^5 n^2 \log^3(n)$$

Exercise 1.5. Define appropriate encodings of the following countable sets into the set of $\{0, 1\}^*$ binary strings:

- The set \mathbb{Q} of all rational numbers.
- The disjoint union $\mathbb{N} \uplus \mathbb{Z}$ of the set \mathbb{N} of the natural numbers and of the set \mathbb{Z} of the integer numbers.
- The class of all finite, directed graphs, namely pairs of the form (V, E) where V is a finite set, and E is a subset of $V \times V$.

2 The Computational Model

Exercise 2.1. Define an efficient 1-tape Turing Machine computing the function $inverse : \{0, 1\}^* \rightarrow \{0, 1\}^*$ such that $inverse(x)$ is the binary string obtained by flipping all bits in x , e.g. $inverse(01011)$ is 10100. Give the Turing Machine explicitly as a triple in the form (Γ, Q, δ) .

Exercise 2.2. Define an efficient 1-tape Turing Machine computing the successor function on the natural numbers, when natural numbers are encoded in pure binary. In order to make your task simpler, you can safely suppose that the binary string encoding a natural number has least significant bits on the left and most significant bits on the right, e.g. 12 is encoded as 0011 rather than as 1100. Give the Turing Machine explicitly as a triple in the form (Γ, Q, δ) .

Exercise 2.3. Do Exercise 2.2, but assume, now, that natural numbers are encoded as usual, e.g., 12 is encoded as 1100.

Exercise 2.4. A pair (x, y) of binary strings *of equal length* can be easily encoded into a single binary string in many ways. Pick one, and write $\sqcup(x, y)\sqcup$ for the encoding of the pair. Define an efficient 3-tape Turing Machine computing the function $xor : \{0, 1\}^* \rightarrow \{0, 1\}^*$ such that $xor(\sqcup(x, y)\sqcup) = x \oplus y$, where \oplus is the bitwise exclusive or operator, i.e. $100 \oplus 101$ is 001. Try to give the Turing Machine explicitly as a triple in the form (Γ, Q, δ) . What happens if we just allow two tapes rather than three?

Exercise 2.5. Show that the function $T(n) = 3n + 2n + 1$ is time-constructible.

Exercise 2.6. Give an example of a function T which is *not* time-constructible, and prove your claim.

3 Polynomial Time Computable Problems

Exercise 3.1. Show that the following problem is computable in polynomial time.

Given a list $A = [a_1, \dots, a_n]$ of natural numbers and a number $v \in \mathbb{N}$, return an index $i \in \{1, \dots, n\}$ such that¹ $A[i] = v$, if any, and return -1 otherwise.

Exercise 3.2. Show that the following problem is computable in polynomial time.

Sort a list $A = [a_1, \dots, a_n]$ of natural numbers.

Hint. You do not need to be efficient: a naive sorting algorithm works fine for solving this exercise.

Exercise 3.3. Determine whether the following algorithms run in polynomial time, where for strings s_1, s_2 we denote by $s_1 :: s_2$ their concatenation. Motivate your answer.

Data: A string $s \in \{0, 1\}^*$

```
p ← s;
ℓ ← |s|;
i ← 1;
while i < ℓ do
  | p = p :: p
end
return p
```

Data: A string $s \in \{0, 1\}^*$

```
p ← s;
ℓ ← |s|;
i ← 1;
while i < ℓ do
  | p = p :: s
end
return p
```

Exercise 3.4. Recall that a *directed* graph is a pair $G = (V, E)$ where V is a set of vertexes and $E \subseteq V \times V$ is the ‘edge’ relation between vertexes. Notice that we do not require E to be symmetric. We represent a graph $G = (V, E)$ with n vertexes $\{v_1, \dots, v_n\}$ using its adjacency matrix, that is the $n \times n$ -matrix A^G defined by $A_{ij}^G = 1$ if $(v_i, v_j) \in E$ and $A_{ij}^G = 0$, otherwise. A universal sink is a vertex v_i such that for all $j, l \leq n$ with $j \neq i$ we have

$$A_{il}^G = 0 \quad A_{ji}^G = 1$$

Notice that if a graph has a universal sink, then the latter is unique. Show that determining whether a graph has a universal sink is computable in polynomial time.

4 Between the Feasible and the Unfeasible

¹We denote by $A[i]$ the i -th element of A .

Solutions to Selected Exercises

Exercise 1.3. For every n , let E_n and O_n be the subsets of $\{0, 1\}^*$ consisting of the strings of length n having even and odd parity, respectively. As an example:

$$\begin{aligned} E_3 &= \{000, 011, 110, 101\}; \\ O_3 &= \{001, 010, 100, 111\}. \end{aligned}$$

It seems that half of the strings of $\{0, 1\}^*$ are in E_3 and half are in O_3 . Is this a general rule. The answer is positive, and indeed, we will now prove that $|E_n| = |O_n| = 2^{n-1}$ for every $n \geq 1$. The first thing we prove is that for every such n , it holds that

$$\begin{aligned} E_n &= \{x \in \{0, 1\}^* \mid x = 0 \cdot y \wedge y \in E_{n-1}\} \cup \{x \in \{0, 1\}^* \mid x = 1 \cdot y \wedge y \in O_{n-1}\} \\ O_n &= \{x \in \{0, 1\}^* \mid x = 0 \cdot y \wedge y \in O_{n-1}\} \cup \{x \in \{0, 1\}^* \mid x = 1 \cdot y \wedge y \in E_{n-1}\} \end{aligned}$$

Every string in E_n , if $n \geq 1$ either starts with a 0 or with a 1. In the former case, the rest of the string is itself in E_{n-1} ; in the latter case, it is of course in O_{n-1} . Viceversa, any string in the form $0 \cdot y$ where $y \in E_{n-1}$ and any string in the form $1 \cdot y$, where $y \in O_{n-1}$ are in E_n . Similarly for strings in the form $0 \cdot y$ where $y \in O_{n-1}$ and any string in the form $1 \cdot y$, where $y \in E_{n-1}$ which are in O_n by definition. As a consequence, we can safely conclude that

$$|E_n| = |E_{n-1}| + |O_{n-1}|; \quad |O_n| = |O_{n-1}| + |E_{n-1}|. \quad (1)$$

The fact that $|E_n| = |O_n| = 2^{n-1}$ for every $n \geq 1$ can thus be proved by induction on n :

- If $n = 1$, then $E_n = \{0\}$ and $O_n = \{1\}$ and the thesis holds.
- Suppose the thesis holds for n , and let us prove that it must hold for $n + 1$:

$$\begin{aligned} |E_{n+1}| &= |E_n| + |O_n| = 2^{n-1} + 2^{n-1} = 2 \cdot 2^{n-1} = 2^n; \\ |O_{n+1}| &= |O_n| + |E_n| = 2^{n-1} + 2^{n-1} = 2 \cdot 2^{n-1} = 2^n. \end{aligned}$$

In both cases, we make use of Equation (1), followed by the induction hypotheses. Please notice how we had to prove a *stronger* result rather than the one required by the exercises, namely that *both* E_n and O_n have a given cardinality. \square

Exercise 1.4. Let us consider, e.g., the functions f_4 and g_4 . Recall that they are defined as follows:

$$f_4(n) = 100n \quad g_4(n) = \frac{1}{100}n \log(n)$$

Intuitively, we expect g_4 to grow asymptotically strictly faster than f_4 . Let us thus prove that f_4 is $O(g_4)$:

$$\begin{aligned} f_4(n) \leq c \cdot g_4(n) &\Leftrightarrow 100n \leq \frac{c}{100}n \log(n) \\ &\Leftrightarrow 100 \leq \frac{c}{100} \log(n) \Leftrightarrow \log(n) \geq \frac{10000}{c} \end{aligned}$$

As a consequence, for every $c > 0$, any n such that $\log(n) \geq \frac{10000}{c}$ would make the inequality $f_4(n) \leq c \cdot g_4(n)$ true, which thus holds for sufficiently large n . Similarly, one can prove that $g_4 \in \Omega(f_4)$:

$$g_4(n) \geq c \cdot f_4(n) \Leftrightarrow \frac{1}{100}n \log(n) \geq 100cn \Leftrightarrow \log(n) \geq 10000c$$

As a consequence, for every c , the inequality $g_4(n) \geq c \cdot f_4(n)$ holds whenever $\log(n) \geq 10000c$, thus for sufficiently large n . \square

Exercise 2.1. The alphabet Γ can be defined as $\{\triangleright, \square, 0, 1\}$, while the set of states Q is $\{q_{\text{init}}, q_s, q_r\}$. The transition function is specified as follows:

$$\begin{aligned}
(q_{\text{init}}, \triangleright) &\mapsto (q_s, \triangleright, S) \\
(q_s, \triangleright) &\mapsto (q_2, \triangleright, R) \\
(q_s, 0) &\mapsto (q_s, 1, R) \\
(q_s, 1) &\mapsto (q_s, 0, R) \\
(q_s, \square) &\mapsto (q_r, \square, S) \\
(q_r, \square) &\mapsto (q_r, \square, L) \\
(q_r, 0) &\mapsto (q_r, 0, L) \\
(q_r, 1) &\mapsto (q_r, 1, L) \\
(q_r, \triangleright) &\mapsto (q_{\text{halt}}, \triangleright, S)
\end{aligned}$$

In all the other cases (e.g. when the state is q_{init} and the symbol is not \triangleright , the behavior of the machine is not relevant, i.e., δ can be defined arbitrarily defined. \square

Exercise 3.1. First, we design an algorithm solving the desired task.

```

Data:  $A = [a_1, \dots, a_n], v$ 
Result: An index  $i \in \{1, \dots, n\}$  s.t.
            $A[i] = v$ , if any,  $-1$ , otherwise
 $i \leftarrow 1$ ;
while  $i \leq n$  do
    if  $A[i] = v$  then
        | return  $i$ 
    else
        |  $i \leftarrow i + 1$ 
    end
end
return  $-1$ 

```

How to prove that this algorithm works in polynomial time? As seen in class, we do that in four steps.

1. We encode the input as a binary string. Our analysis of the complexity of the algorithm will be given with respect to the length (call it ℓ) of such a string.
2. We prove that the number of instructions of the algorithm is bounded by a polynomial in ℓ .
3. We argue that each instruction can be simulated by a Turing machine in polynomial time.
4. We show that all ‘intermediate’ data and results of the algorithm are bounded by a polynomial in ℓ .

We begin with step 1. In order to encode $A = [a_1, \dots, a_n]$ as a binary string, we first need to encode its components a_i . For that, we can choose one standard encoding of the natural numbers in $\{0, 1\}^*$. Let us write $\lfloor a_i \rfloor$ for the encoding of a_i in binary. Since using n bits we can encode $2^n - 1$ (natural) numbers, the encoding of a number $a \in \mathbb{N}$ requires $\log a + 1$ bits². Therefore, we have $|\lfloor a_i \rfloor| = \log a_i + 1$. The next step is to understand how to encode the whole A . For that, we regard a list of elements $[b_1, \dots, b_n]$ as a ‘pair of pairs’ of the form $((b_1, b_2), b_3), \dots, b_n$. Recall that given a pair (b_1, b_2) of bitstrings, we can define the string $\lfloor (b_1, b_2) \rfloor \in \{0, 1\}^*$ by first translating (b_1, b_2) as the string $b_1 \# b_2 \in \{0, 1, \#\}$ and then encoding $b_1 \# b_2$ as a string $\lfloor (b_1, b_2) \rfloor \in \{0, 1\}^*$. For the latter point, we simply map 0 to 00, 1 to 11, and $\#$ to 01. As a consequence, we see that $|\lfloor (b_1, b_2) \rfloor| = 2|b_1| + 2|b_2| + 2$. Now, given a list $[b_1, \dots, b_n]$ of bitstring by regarding it as a pair

²Actually, we should take the floor of $\log a$.

$((b_1, b_2), b_3), \dots, b_n)$, we see that we obtain an encoding $\sqcup[b_1, \dots, b_n]_{\sqcup}$ of length $\sum_{i=1}^n 2|b_i| + 2(n-1)$. Applying these general considerations to A (and recalling that $|\sqcup a_i \sqcup| = \log a_i + 1$), we obtain:

$$\sqcup A \sqcup = \sqcup[\sqcup a_1 \sqcup, \dots, \sqcup a_n \sqcup]_{\sqcup} \qquad |\sqcup A \sqcup| = \sum_{i=1}^n 2(\log a_i + 1) + 2(n-1)$$

Finally, we pair A with v (recall that both A and v are input of our algorithm), so $\sqcup(\sqcup A \sqcup, \sqcup v \sqcup)_{\sqcup}$ gives an encoding of the input of our algorithm in binary notation. Notice that

$$\ell = |\sqcup(\sqcup A \sqcup, \sqcup v \sqcup)_{\sqcup}| = 2\left(\sum_{i=1}^n 2(\log a_i + 1) + 2(n-1)\right) + 2(\log v + 1) + 2.$$

We now move to step 2. The latter is straightforward. Our algorithm consists of:

- 1 assignment ($i \leftarrow 1$).
- n iteration of:
 - An inequality check ($i \leq n$).
 - A conditional branching performing:
 - * An equality check ($A[i] = v$),
 - * Either a return instruction or an assignment ($i \leftarrow i + 1$).
- 1 return instruction

Therefore, the number of the instruction is of the form $b + c \cdot n$, for suitable constants b, c , and thus it is bounded by a polynomial in ℓ . In order to prove step 3, we have to argue that all the aforementioned instructions can be simulated by a TM in polynomial time. For instance, an equality check can be simulated as follows. Say we have two values a and b stored in different portions of a tape of a TM. In order to check whether a is equal to b , the machine simply moves back and forth between a and b checking whether they are bitwise equal. This can be done in polynomial time with respect to the length of a and b , provided that the ‘distance’ between a and b in the tape is itself bounded by a polynomial in the length of a and b . This will be indeed ensured by step 4. Similar arguments can be used to show that all other instructions can be simulated efficiently by a TM.

Finally, in order to prove step 4 we simply observe that the only intermediate value computed by our algorithm is i , which can be at most n (and therefore it is bounded by a polynomial in ℓ). \square

Exercise 3.3. In order to prove that the two algorithms run in polynomial time we have to follow the four steps of previous exercise. These are mostly straightforward. In fact, step 1 is trivial, as the input s is already in binary. For step 2 we simply observe that we have ℓ iterations, and that overall the number of instructions is of the form $b + c \cdot \ell$, for suitable constants b, c , and thus polynomial in ℓ . Moreover, it is not hard to see that all the instructions can be easily simulated by a TM. What goes wrong is step 4. In fact, in the first algorithm at each iteration we concatenate p with itself. That means that if before entering into the while-loop $|p| = \ell$, then after one iteration we will have

$$|p| \text{ after iteration 1} = 2(|p| \text{ at iteration 0}) = 2\ell$$

Similarly, we obtain

$$\begin{aligned} |p| \text{ after iteration 2} &= 2|p| \text{ (at iteration 1)} = 4\ell \\ |p| \text{ after iteration 3} &= 2|p| \text{ (at iteration 2)} = 8\ell \\ &\vdots \end{aligned}$$

and thus

$$|p| \text{ after iteration } n = 2(|p| \text{ at iteration } n - 1) = 2^n \ell$$

That means that the length of p is exponential in ℓ , and therefore the first algorithm cannot run in polynomial time. Notice that this is not true for the second algorithm, where we have $|p| \approx \ell^2$.
 \square