

Non Monotonic Logics



Prof. Mauro Gaspari
Dipartimento di Informatica Scienze e Ingegneria
(DISI)

mauro.gaspari@unibo.it

Problems with FOL



- The size of FOL KB necessary to describe real situations is almost always too large.
- Often available knowledge is incomplete and FOL has not specific mechanisms to deal with incomplete knowledge. For example FOL is too rigid to deal with new (conflicting) knowledge.
- To model commonsense reasoning, it is necessary to be able to jump to plausible conclusions from the given knowledge making assumptions.
- The choice of assumptions is not blind: most of the knowledge on the world is given by means of general rules which specify typical properties of objects. For instance, "birds fly" means: birds typically fly, but there can be exceptions such as penguins, ostriches, ...

Example



$\forall x. \text{Bird}(x) \rightarrow \text{Fly}(x)$




FOL solution



- We can try to capture this semantics adding more conditions:

$$\forall x \text{ Bird}(x) \wedge \neg \text{Penguin}(x) \wedge \neg \text{Ostrich}(x) \wedge \neg \text{Peacock}(x) \wedge \neg \text{Chicken}(x) \rightarrow \text{Fly}(x)$$

- But this is not enough: indeed in general situations additional exceptions should be considered.


$$\forall x \text{ Bird}(x) \wedge \neg \text{Penguin}(x) \wedge \neg \text{Ostrich}(x) \wedge \neg \text{Peacock}(x) \wedge \neg \text{Chicken}(x) \wedge \neg \text{Tied}(x) \wedge \neg \text{Caged}(x) \wedge \neg \text{Cemented}(x) \rightarrow \text{Fly}(x)$$

And More



- $\forall x \text{ Bird}(x) \wedge \neg \text{Penguin}(x) \wedge \neg \text{Ostrich}(x) \wedge \neg \text{Peacock}(x) \wedge \neg \text{Chicken}(x) \wedge \neg \text{Tied}(x) \wedge \neg \text{Caged}(x) \wedge \neg \text{Cemented}(x) \wedge \neg \text{Injured}(x) \wedge \neg \text{Dead}(x) \wedge \neg \text{Baby}(x) \rightarrow \text{Fly}(x)$

Observation



- It is always possible to add more detail to the environment but the Frame Problem arises.
- In general it is not possible to know which are the right exceptions to add. It is possible only if the possible queries are known.



Nonmonotonic Reasoning



- **Nonmonotonic reasoning** deals with the problem of deriving plausible conclusions, but not infallible, from a KB.
- Since the conclusions are not certain, it must be possible to retract some of them if new information shows that they are wrong.
- Classical logic is inadequate since it is monotonic: if a formula B is derivable from a set of formulas S , then B is also derivable from any superset of S :

$S \vdash B$ implies $S \cup \{A\} \vdash B$, for any formula A .

- Nonmonotonic logics have been devised with modified notions of truth and entailment in order to capture such behavior.

Example



- Let the KB contains:

Typically birds fly.

Penguins do not fly.

Tweety is a bird.


- We can infer that Tweety flies.
- However if the following information is added to the KB:

Tweety is a penguin

- the previous conclusion must be retracted and, instead, the new conclusion that Tweety does not fly will hold.

Nonmonotonic Reasoning



- The statement "typically A" can be read as: "in the absence of information to the contrary, assume A". 
- The problem is to define the precise meaning of "in the absence of information to the contrary".
- The meaning could be: "there is nothing in KB that is inconsistent with assumption A".
- However, other interpretations are possible
- Different interpretations give rise to different non-monotonic logics


Nonmonotonic logics



- Non-Monotonic logics have been proposed from the beginning of the 80's, here are historically the most important proposals:
 - **Closed World Assumption.**
 - **Circumscription, by McCarthy, '80**
 - Non-monotonic logic, by McDermott and Doyle, '80
 - **Default Logic, by Reiter, '80**
 - Autoepistemic logic, Moore '84

Closed World Assumption



- A basic understanding of the logic used in databases, is that only positive information is represented explicitly and negative information is not represented explicitly.
- A more sophisticated KR system might allow the user to specify rules for which to apply the Closed World Assumption (CWA). This logic can be also applied to Kbs. If a positive fact is not present in a KB, it is assumed that its negation holds.
- This is called Closed World Assumption: the only true facts are the provable ones, this is the principle of negation as (finite) failure.
- If $KB \not\models A$ then ~~DB~~ $\vdash_{CWA} \neg A$ 
- Note that, this inference is not valid in classical logic.

Example



- Suppose a KB contains facts of the form "practice(person, sport)":

practice(anne, tennis)

practice(joe, tennis)

ractice(anne, sky)

- Then we have:

~~DB~~ $\vdash_{CWA} \neg \text{practice}(\text{joe}, \text{sky})$

- Trivially CWA is non-monotonic, since adding a fact may lead to withdraw the negative conclusion:

- ~~DB~~ $\cup \{\text{practice}(\text{joe}, \text{sky})\} \not\vdash_{CWA} \neg \text{practice}(\text{joe}, \text{sky})$

Circumscription



- **Circumscription** can be seen as a more powerful and precise version of the closed world assumption.
- The idea is to specify particular predicates that are “as false as possible”: false for every object except those for which they are known to be true.
- For example, suppose we want to assert the default rule that birds fly. This can be done introducing a predicate $\text{Abnormal}_1(x)$:
$$\forall x \text{ Bird}(x) \wedge \neg \text{Abnormal}_1(x) \rightarrow \text{Flies}(x)$$
- If we specify that Abnormal_1 is to be circumscribed a circumscriptive engine assumes $\neg \text{Abnormal}_1(x)$ unless $\text{Abnormal}_1(x)$ is known to be true.
- We can infer $\text{Flies}(\text{Tweety})$ from $\text{bird}(\text{Tweety})$, but the conclusion no longer holds if $\text{Abnormal}_1(\text{Tweety})$ is asserted.

Model Preference Logics




- Circumscription can be viewed as an example of **model preference logic**.
- In these logics a sentence is entailed (as a default) if it is true in all preferred models of the KB (in classical FOL it should be true in all models).
- A model is preferred to another if it has fewer abnormal objects.

Example: Nixon Diamond



- This approach can be exploited to model multiple inheritance in the context of semantic networks.
- The Nixon diamond: Is Nixon a Pacifist?



Have you neve been
In Vietnam?

Example



- We can encode this knowledge in the following KB:

$\text{Republican}(\text{Nixon}) \wedge \text{Quaker}(\text{Nixon})$

$\text{Republican}(x) \wedge \neg \text{Abnormal}_2(x) \rightarrow \neg \text{Pacifist}(x)$

$\text{Quaker}(x) \wedge \neg \text{Abnormal}_3(x) \rightarrow \text{Pacifist}(x)$

- If we circumscribe Abnormal_2 and Abnormal_3 there are two preferred models:
 1. $\text{Abnormal}_2(\text{Nixon}), \neg \text{Pacifist}(\text{Nixon})$
 2. $\text{Abnormal}_3(\text{Nixon}), \text{Pacifist}(\text{Nixon})$
- Thus the circumscriptive engine remains properly agnostic.
- If we want to assert that religion beliefs take precedence over political beliefs we can use a formalism called **prioritized circumpcription**.

Preferential Models



- Define a pair $\langle L, \sqsubseteq \rangle$ where L is a non monotonic logic (based on FOL) and \sqsubseteq is a partial order on the models.
- $M_1 \sqsubseteq M_2$ reads as M_2 is preferred over M_1 .
- If M is a model of KB , M is a preferred model if no other model M' is a model of KB and $M \sqsubseteq M'$.
- We say that $KB \models_{\sqsubseteq} P$ iff P is true for all preferred models of KB .

Default Logic



- Default logic extends FOL inference rules by non-standard **default rules**. These rules allows one to express default properties.
- Example:

$$\frac{\text{bird}(x) : \text{flies}(x)}{\text{flies}(x)}$$

- This rule can be interpreted as: "if x is a bird and we can consistently assume that x flies (flies(x) is consistent with the KB) then we can infer that x flies".

Default Rules



- In general a Default Rule has the form:

$$\frac{\alpha(x) : \beta(x)}{\gamma(x)}$$

- that can be interpreted as: "if $\alpha(x)$ holds and $\beta(x)$ can be consistently assumed then we can conclude $\gamma(x)$ ".
 - $\alpha(x)$: the prerequisite
 - $\beta(x)$: the justification
 - $\gamma(x)$: the consequent
- It is possible to specify more than one justification: in any one of them can be proven false, then the conclusion cannot be drawn.

Default theory



- A default theory is a pair $\langle D, W \rangle$, where D is a set of default rules and W is a set of first-order formulas.
- Example 1: let $\langle D, W \rangle$ be:

$$D = \left\{ \frac{bird(x) : fly(x)}{fly(x)} \right\}$$

$$W = \{bird(tweety), \forall x(penguin(x) \rightarrow bird(x)), \\ \forall x(penguin(x) \rightarrow \neg fly(x))\}$$

Extension



- Intuitively, in a default theory $\langle D, W \rangle$: W represents the stable (but incomplete) knowledge of the world
- D rules for extending the knowledge W by plausible (but defeasible) conclusions.
- Notion of **extension** of a default theory: the theory (= deductively closed set of logical formulas) obtained by extending W by the rules in D .

Example 2



- Let $\langle D, W \rangle$ be as in the previous example 1.
- Since $\text{bird}(\text{tweety})$ is true, and it is consistent to assume $\text{fly}(\text{tweety})$, then $\text{fly}(\text{tweety})$ is true in the (unique) extension of $\langle D, W \rangle$.
- Consider now the the default theory $\langle D, W' \rangle$, where $W' = W \cup \{\text{penguin}(\text{tweety})\}$
- then the assumption $\text{fly}(\text{tweety})$ is no longer consistent, and the application of the default rule is blocked.

Example 3



- Let $\langle D, W \rangle$ be as follows:

$$D = \left\{ d_1 = \frac{Rep(x) : \neg Pac(x)}{\neg Pac(x)}, d_2 = \frac{Quack(x) : Pac(x)}{Pac(x)} \right\}$$

$$W = \{Rep(Nixon), Quack(Nixon)\}$$

- For both default rules d_i , the prerequisite is derivable from W . What can be concluded from $\langle D, W \rangle$?
- If we apply d_1 , we conclude $\neg Pac(Nixon)$; therefore $Pac(Nixon)$ cannot be assumed consistently, so that d_2 is blocked.
- If we apply d_2 , we conclude $Pac(Nixon)$; therefore $\neg Pac(Nixon)$ cannot be assumed consistently, so that d_1 is blocked.

Observations



- There are two extensions: one containing $\neg \text{Pac}(\text{Nixon})$ and the other containing $\text{Pac}(\text{Nixon})$.
- An extension (to be defined next) represents the set of plausible conclusions.
- In summary a default-theory may have zero, one, or many extensions.
- How to deal with this issue?

Inference



- Since a default theory $\Delta = \langle W, D \rangle$ may have multiple extensions (including none), how to define a notion of inference?
- There are two natural notions:
 - **Credulous inference:** $\Delta \vdash_c A$ if there exists an extension E of Δ such that $A \in E$.
 - **Skeptical inference:** $\Delta \vdash_s A$ if for all extensions E of Δ , we have $A \in E$.
- Note that, since a default theory may have no extensions $\Delta \vdash_s A$ does not imply $\Delta \vdash_c A$.

Truth Maintenance Systems



- We have argued that many of the inferences drawn from a KRS will have only default status rather than being absolutely certain.
- **Belief revision** is sometimes necessary: when some of the inferred facts turn out to be wrong.
- A **truth maintenance system** (TMS) is a program that keeps track of dependencies between sentences so that retraction (and some other operations) will be more efficient.

A Simple TMS



- The simplest approach to truth maintenance is to use **chronological backtracking**.
- In this approach, we keep track of the order in which sentences are added to the knowledge base by numbering them from P_1 to P_n . When the call $\text{RETRACT}(\text{KB}, P_i)$ is made, the system reverts to the state just before P_i was added. Removing P_i and the sentence derived from it.
- If desired, the sentences P_{i+1} through P_n can then be added again. This is simple, and it guarantees that the knowledge base will be consistent, but it means that retraction is $O(n)$, where n is the size of the knowledge base.
- We would prefer a more efficient approach that does not require us to duplicate all the work for P_{i+1} to P_n .

TMS Functionalities



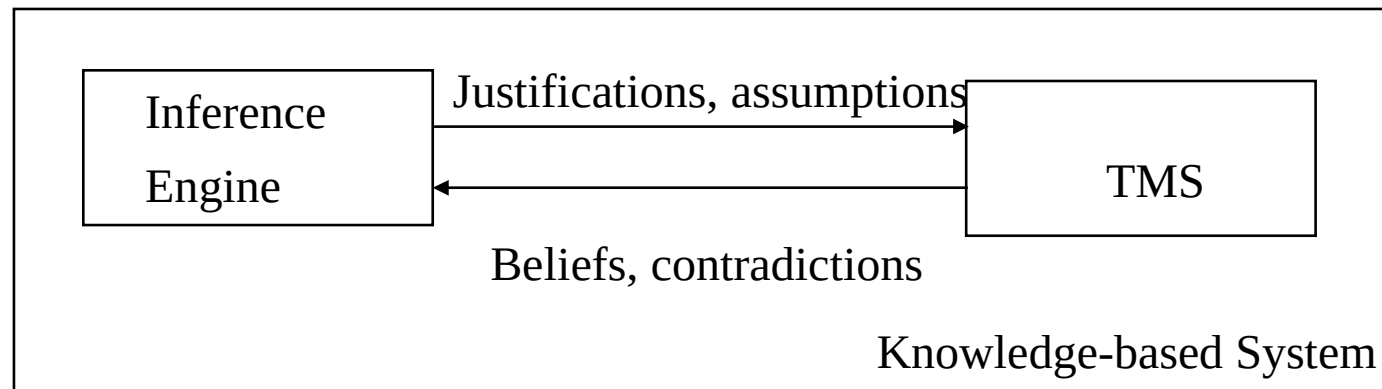
- A TMS actually performs four important jobs:
 1. It enables dependency-directed backtracking, to avoid the inefficiency of chronological backtracking.
 2. It provides explanations of propositions. A proof is one kind of explanation—if we ask, "Explain why you believe P is true?" then a proof of P is a good explanation. If a proof is not possible, then a good explanation is one that involves assumptions.
 3. It support default reasoning.
 4. It helps in dealing with inconsistencies. If adding P to the KB results in a logical contradiction, a TMS can help pinpoint an explanation of what the contradiction is.



TMS



- A Truth Maintenance System (TMS) is module of a KRS that can be integrated with an engine providing the described functionalities.



TMS Types



- There are several types of TMSs:
- **JTMS** (*justification-based TMS*): is the simplest system. In a JTMS, each sentence in the knowledge base is annotated with a justification that identifies the sentences from which it was inferred, if any. For example, if Q is inferred by Modus Ponens from P , then the set of sentences $\{P, P \rightarrow Q\}$ could serve as a justification of the sentence Q .
- **ATMS** (*assumption-based TMS*): it is more powerful and maintains information at each node within a single dependency network for all the alternative sets of assumptions that we may hold, also called worlds or environments.
- **LTMS** (*logic-based TMS*): Nodes have three possible labels ($:TRUE$, $:FALSE$, $:UNKNOWN$) and Justifications are disjunctive clauses.

Enforcing logical relations among beliefs



- Most of AI problem solvers are based on search
- Changing assumptions requires updating consequences of beliefs
- Rederivation can be expensive
 - Large, complex calculations (e.g., computational fluid dynamics)
 - Physical experiments

Recover from Inconsistencies



- Data can be wrong
 - The patient's temperature is 986 degrees.
- Constraints can be impossible
 - Our new computer should
 - »Run off batteries for 8 hours
 - »Fit in a shirt pocket
 - »Run faster than a Cray MP-X



Explanations

- Providing answers is not enough
 - Cut the patient's heart out
 - That design won't work
- Explanations are needed
 - Radical bypass surgery is required because...
 - No material will stand the projected stresses.



Guide Backtracking

- (Dependency-directed backtracking vs Chronological backtracking).
- Avoid rediscovering contradictions
- Avoid throwing away useful results



Example

Choose in sequence:

- A or B
- C or D
- E or F

Given: A and C cannot hold together


Given: B and E cannot hold together

Assume we want all consistent solutions

Assume that we cannot test until every choice has been made




Search Space



Chronological Backtracking




- Often wastes computation: Suppose that D and F together cause lots of work. Popping context loses this work



Chronological Backtracking




Rediscovered Contradictions: we can avoid to try B and E together more than once.



Guide backtracking



- Dependencies can guide backtracking:





JTMS Idea

- *Justifications* express relationships between beliefs
 - Justifications for a belief provide explanations, ability to pinpoint culprits
- Belief in an assertion expressed via its *label*
 - P being in database no longer is the same as believing P
 - Assertions and justifications serve as cache
 - Rules/other computations need only be executed once
- Justifications can be used to record inconsistencies
 - *Dependency-directed backtracking*
- Defaults can be represented via explicit assumptions



JTMS nodes

- Each belief is represented by a TMS node
- TMS nodes are associated 1:1 with assertions
- The label of a node represents the belief status of the corresponding problem solver fact.
- The relationships between beliefs are expressed by the justifications it participates in.



JTMS Labels

- Every assertion is either IN or OUT
 - IN = “believed”
 - OUT = “not believed”
- Warning: IN does not mean TRUE




	P in	P out
(not P) in	Contradiction	(not P) true
(not P) out	P true	Don't know

JTMS Justifications




- Must be Horn clauses
- Definitions
 - *Consequent* is the node whose belief is supported by a the justification
 - *Antecedents* are the beliefs which, when IN, support the consequent
 - *Informant* records information from external systems



Dependency Networks



- Each node has:
 - *Justifications*: the justifications which have it as the consequent
 - *Consequences*: justifications which use it as an antecedent
 - *Support*: a single justification taken as the reason for it being IN, if any.




Special states of JTMS nodes




- *Assumptions* are IN if enabled.
- *Premises* are always IN.
- *Contradictions* should never hold.



Assumption



Premise



Contradiction



■ Enforcing constraints between beliefs

- A node is IN when either:
 1. It is an enabled assumption or premise
 2. There exists a justification for it whose antecedents are all IN
- Assumptions underlying a belief can be found by backchaining through supporting justifications
- JTMS operations must preserve *well-founded support*.



A TMS operates incrementally


- At any time, the inference engine can add
 - new justifications
 - declare a statement to be a premise or contradiction (permanently)
- Assume a statement
- In all cases,
 1. Set the directly affected node, if any.
 2. Propagate the consequences (Propagation)



Example

Propagation of Belief


Initial state of dependency network:





Example


Suppose inference engine enables A:





Example


D becomes believed via J1:





Example

F becomes believed via J3:



Retracting information




- Premises, contradictions cannot be retracted
- Justifications cannot be retracted
 - They comprise the problem solver's cache
 - Rules need only be run once for each set of matching data
- Assumptions can be retracted
- Algorithm:
 - 1 Make assumption OUT
 - 2 Retract all nodes which rely on it (Propagate-outness)
 - 3 Find alternate support for newly OUT nodes.



Retraction Example


Initial state:





Retraction Example


Retract C:





Retraction Example


E becomes out:





Retraction Example


F becomes out:



Retract then Resupport




Initial state:



Retract then Resupport




Retract A:



Retract then Resupport




Retract D via J1:



Retract then Resupport



Resupport D via J4:



Loss of well-founded support!

References



- R. Reiter (1980). A logic for default reasoning. *Artificial Intelligence*, 13:81-132.
- Stuart C. Shapiro, Belief Revision and TMSs: An Overview and a Proposal CSE Technical Report 98-10.
- J. Doyle, A truth maintenance system. *Artificial Intelligence*, Volume 12, Issue 3, 1979, Pages 231-272.
- Johan de Kleer, An assumption-based truth maintenance system *Artificial Intelligence*, Volume 28, Issue 2, 1986, Pages 127-162.

TOOLS



- Truth-teller: An assumption based truth maintenance system (ATMS) in C#:

<https://github.com/maate/truth-teller>

- JTMS: A simple implementation of Doyle's Justification-based Truth Maintenance System (JTMS):

<https://github.com/hbeck/jtms>

- ATMS-in-Python: Assumption-based truth maintenance systems implementation in Python 3 accordingly to the work of Johan de Kleer:

<https://github.com/FellnerDotDev/ATMS-in-Python>

The End



Byebye!
See you at the exam!
I will be there!